

Bits / Bytes	Audio storage	Paging formulas	Calculation	IP address counts	Bit shifts	Pipelining
• 1 byte = 8 bits	Number of samples:	Let page size $P$ (in bytes). $N_{samples} = f_s \times T$ where $f_s$ = sample rate (Hz), $T$ = time (s)	Split virtual address: $\text{page\_number} = \frac{ VA }{P}$ $\text{offset} = V \bmod P$	IPv4: #addresses = $2^{32}$ IPv6: #addresses = $2^{128}$	Let $xbe an integer.$ Left shift by $a$ bits: $x \ll a = x \times 2^a$ Right shift by $a$ bits (integer): $x \gg a \approx \frac{ x }{2^a}$	Let stages be $t_a, t_b, t_c, t_d$ . Non-pipelined instruction time: $t_{inst} = t_a + t_b + t_c + t_d$
• 1 KB $\approx 2^{10}$ bytes = 1024 B	Total bits:	bits = $N_{samples} \times \text{bits per sample}$	Physical address (given page frame number PFN): $PA = PFN \times P + offset$	Integer ranges	• Unsigned n-bit: $0 \leq x \leq 2^n - 1$	Non-pipelined throughput (instructions per second): $\frac{1}{t_{inst}}$
• 1 MB $\approx 2^{20}$ bytes	Total bytes:	bytes = $\frac{\text{bits}}{8}$	Number of virtual pages: $\#pages = \frac{VM \text{ size}}{P}$	Signed 2's complement (n bits): $-2^{n-1} \leq x \leq 2^{n-1} - 1$	• Pipeline stage time: $t_{stage}$	
• 1 GB $\approx 2^{30}$ bytes	Number of values / addresses	Values from n bits: $\#values = 2^n$	Number of page frames in RAM: $\#frames = \frac{RAM \text{ size}}{P}$	Cache performance	• Hit ratio: $H = \frac{\text{hits}}{\text{hits} + \text{misses}}$	Pipelined throughput (ideal): $IPS_{pipe} = \frac{1}{t_{stage}}$
<b>Base + Limit (simple virtual memory)</b>	Physical address:	If valid: $PA = BASE + VA$	Effective access time (EAT): $EAT = H \cdot t_{cache} + (1 - H) \cdot t_{mem}$	• Pipeline speed-up (ideal): $\text{Speed-up} \approx \frac{t_{inst}}{t_{stage}}$	Pipeline stage time: $t_{stage}$	

## 1. OPERATING SYSTEMS

→ **Basic picture**

User ↔ Application ↔ Operating System (OS) ↔ Hardware.

OS = special software that is loaded by boot code when the computer powers on.

→ **What is an OS**

Sequence of instructions that controls the machine hardware and provides an environment for applications.

Most users never see it directly – but you need to.

→ **OS taxonomy / types**

**Single-user** (e.g. typical desktop), **Multi-user** (shared servers, mainframes), **Embedded** (routers, phones, IoT devices), **Cluster / distributed systems** (many machines cooperating).

Real-time OS → CPU, memory, devices, files. | **Virtual machine concept** – each program thinks it owns the machine. | **Implements security & protection**

polices – who can access what.

→ **OS structure models**

**Monolithic** – big blob of kernel code that does (almost) everything. | **Layered** – OS split into layers that build on one another. | **Microkernel** – very small kernel (scheduling, IPC, basic memory); other services in user-space.

→ **9 Classic OS features**

Support for concurrency Many tasks "at once" using **context switching**, **scheduling**, protection, synchronisation.

Sharing mechanisms Sharing code/resources like drivers, window system; one copy used by many processes.

Persistent storage Long-term data on discs/SSD/HDD; convenient, reliable access; backups.

Nondeterminacy Events arrive in unpredictable order/timing: OS must still behave correctly.

Efficiency High utilisation, low overhead, fast response.

Reliability Minimise bugs' impact; maximise Mean Time Between Failures (MTBF).

Resilience Tolerate hardware faults where possible.

Maintainability Code easy to fix/extend.

Small size Less memory, less chance of bugs, easier updates.

→ **Processes & concurrency**

OS = many processes: some for system, some for users. | **Process** = running program (code + data + state). | **Multiple processors** (CPU cores, GPUs, FPGAs) can run processes in parallel. | To support concurrency, processes must: **Communicate** – share data. **Synchronise** – coordinate ordering.

→ **Two core concurrency concepts**

**Mutual exclusion (mutex)** More than one process wants a shared but unshareable resource (printer, table row, mailbox). Only one can use it at a time.

**Synchronisation** One process must wait for another to reach a certain point before continuing.

→ **OS Kernel**

**Kernel** = most critical part of OS. Is either: Core layer of layered system; or Main body of a monolithic OS. | Often split into **upper** and **lower** halves; lower is the Hardware Abstraction Layer (HAL) that hides hardware details.

→ **Kernel responsibilities**

Process switching & scheduling, Inter-process communication (IPC) & synchronisation, | **Hardware interfacing** (drivers).

→ **Kernel components**

First-level interrupt handler – reacts instantly to interrupts | Communication primitives – pipes, message queues, semaphores, etc | **Low-level scheduler** – chooses which process/thread runs next.

→ **Process Control Block (PCB)**

1. Data structure used by OS to describe a process.

2. Typically holds ID | state (running/readyblocked) | registers | program counter | stack pointer | memory info (base/limit or page table), scheduling info | open files, etc.

## 2. MEMORY MANAGEMENT, VIRTUAL MEMORY AND PAGES

→ **Basic idea of memory**

1. Stores **binary data**: numbers, text, images, programs. | Programmers think in **variables & types**; hardware sees only **addresses & bits**. | Same data comes back out if you read from the same address.

→ **Real vs virtual addresses**

Real (**physical address**) – location in actual RAM. Unique. | **Virtual address** – what your program sees. | OS + hardware (MMU) translate virtual → physical. | Benefits: isolation, protection, easier programming.

→ **What OS does for addresses**

Address translation – via MMU & page tables / base+limit. | **Address checking** – ensure within legal range. | **Memory protection** – read/write/execute permissions.

→ **2.1 Base + limit (Simple virtual memory)**

→ **Idea**

Early CPUs had simple hardware: each process gets a contiguous region of physical memory.

Two registers per process: **BASE**: starting physical address of process. **LIMIT**: size of its address space.

→ **Address translation**

Program uses **VirtualAddress** (0..LIMIT-1). | **Hardware checks VirtualAddress < LIMIT**. | **PhysicalAddress = BASE + VirtualAddress** if valid, else memory fault.

→ **4 PPs / com**

- Simple, fast hardware.

- Leads to fragmentation, poor utilisation, hard to grow address space.

→ **2.2 Paging systems (modern virtual memory)**

→ **Concepts**

Virtual memory and physical memory are both split into fixed-size blocks: **Pages** (virtual), **Page frames** (physical). MMU translates (page\_number, offset) → (frame\_number, offset).

→ **Address breakdown**

Given page size **P** bytes: | page\_number = floor(VA / P) | offset = VA mod P | RealAddress = frame\_number × P + offset

→ **Page tables**

OS stores, per process, a **page table** with one entry per virtual page. Each Page Table Entry (PTE) typically has: | **V (valid)** bit – is this page currently in a frame? | **PROT / protection** – read/write/exec permissions/users. | **M (modified/dirty)** bit – page changed since loaded. | **OWN** / software bits – used by OS. | **PFN** – page frame number.

→ **Page faults when they occur**

Page number invalid / out of range. | Page beyond allocated virtual memory. | Legal but never referenced before (needs to be created). | Page is currently swapped out to disk (page file). | Access violates permissions.

→ **Page-in operation (handling a "not in RAM yet" fault)**

Find a free page frame (or pick a victim to evict). | Read required page from disk into that frame. | Update PTE (set V bit, PFN, clear M, etc.). | Restart / resume the faulting instruction.

→ **Fetch strategies**

Demand fetch – only load page when first referenced. | **Anticipatory / prefetch** – guess which pages will be needed soon and load early.

→ **Retirement / replacement strategies**

Random – pick any page frame. Simple but poor. | **Least recently used (LRU)** – evict page that hasn't been used for the longest time. Better locality but more overhead.

→ **Overheads**

Address translation requires at least one extra memory access (to read PTE) → effectively "double read" unless you use a cache like TLB.

## 3. FILE SYSTEMS

→ **Profile of file systems**

Long-term information storage – reliable, efficient, long-lasting. | Store data on media like HDD/SSD, optical discs, flash drives, etc. | Provide **security** and **access control**. | Survive reboots and crashes (ideally).

→ **Common file systems**

FAT, NTFS (Windows), HFS/APFS (Apple), ext/3/4, btrfs, xfs, zfs, ufs, vxf, ISO9660 (CD-ROM), NFS, ReiserFS, etc.

→ **What are files?**

Containers for: data | program source | binaries | configuration, etc.

File system must support: I. Creation, deletion, modification. | Read, write, append. | Naming and directories. | Protection, privacy, security. | Efficient access.

→ **Core components**

One gets a handle / descriptor | **Read** – from current offset | **Write** – to current offset | **Seek/rewind** – move read/write pointer | **Close** – release descriptor.

→ **3.1 File organisation on disk**

Contiguous organisation

File stored in a sequence of consecutive blocks on disk. | Directory stores start block and length.

→ **Advantages**

Avoids external fragmentation (any free block can be used). | Easy to grow/shrink files.

→ **Disadvantages**

Map itself can be large and must be read → overhead. | Random access is slow (must follow pointers). | Extra disk accesses to traverse the map.

→ **3.2 Directories & metadata**

File name and type. | Permissions and access control bits. | Owner and group. | Timestamps – creation, last access, last modification, possibly last backup. | File size and/or block count. | Pointer(s) to data blocks / inodes.

→ **3.3 Maintenance & backup**

File system maintenance

Compaction – rearrange files to reduce fragmentation. | **Defragmentation** – move blocks so each file is more contiguous. | **Consistency checks** (fsck, chkdsk)

detect/correct metadata errors. | **Formatting** – create new, empty file system structures. | **Detect bad blocks** – mark them unusable and map around them.

→ **Backup strategies**

Goal: be able to **restore part or all** of the file system after failure.

Types: **Full backup** – everything. | **Differential** – changes since last full backup. | **Incremental** – changes since last any backup (full or incremental).

Requirements: reliable | complete | correct | minimal redundancy | physically secure.

Media: cloud | second internal disk | external HDD/SSD | tape | USB | optical discs.

## Calculation

Page size $P$ (in bytes).	Let $N_{samples} = f_s \times T$	Let page size $P$ (in bytes).	Let $xbe an integer.$
where $f_s$ = sample rate (Hz), $T$ = time (s)		Split virtual address: $\text{page\_number} = \frac{ VA }{P}$	
Total bits:	bits = $N_{samples} \times \text{bits per sample}$	Physical address (given page frame number PFN): $PA = PFN \times P + offset$	
bits = $\frac{bytes}{8}$		Number of virtual pages: $\#pages = \frac{VM \text{ size}}{P}$	
Number of page frames in RAM:	$\#frames = \frac{RAM \text{ size}}{P}$	Cache performance	
#addresses = $2^n$			
		Hit ratio: $H = \frac{\text{hits}}{\text{hits} + \text{misses}}$	
		Effective access time (EAT): $EAT = H \cdot t_{cache} + (1 - H) \cdot t_{mem}$	

## 4. NETWORKING PART 1 – Connectivity, IP Model & Ethernet

→ **What is a computer network?**

System that allows data interchange between computers. | A simple peripheral (e.g. printer directly plugged into a PC) is not usually called a "network".

→ **Physical connectivity**

Links can be wired (Ethernet, fibre, USB, CAN, etc.) or wireless (Wi-Fi, 5G, Bluetooth, etc.). | Network can connect two machines (point-to-point) or many machines (LAN, WAN, Internet).

→ **Internet**

A connection protocol between computers; can be one hop or many hops over heterogeneous physical links.

→ **4.1 OSI model (reference)**

+7-layer OSI model (classic)

Physical layer | Data link layer | Transport | Session | Presentation | Application

(Used mainly as a conceptual model; real Internet stack is simpler.)

→ **4.2 Internet / TCP/IP-5 layer model**

+Layers (bottom → top)

Physical – sends bits across a medium. | **Network Interface** – groups bits into frames/packets for specific media; hides differences between local networks. | Internet layer – IP (Internet Protocol) provides best-effort, end-to-end delivery of datagrams (no guarantee of arrival). | **Transport layer** – splits / reassembles data into segments/datagrams; uses protocols such as TCP and UDP. | **Application layer** – protocols such as HTTP, SMTP, POP3, FTP, NFS, SMB/NETBIOS, etc.

→ **Encapsulation idea**

Each layer wraps the data from above with its own **header** (and sometimes footer). | On receive, headers are stripped layer by layer.

→ **4.3 Transport layer basics (TCP/UDP)**

+ UDP (User Datagram Protocol)

Best effort, connectionless, no guarantee of delivery/order. | Very low overhead, good for time-critical traffic (DNS queries, streaming, VoIP). | Endpoint identified by (IP address, port).

→ **TCP (Transmission Control Protocol)**

Connection-oriented: uses handshake to establish a reliable session. | Guarantees in-order, error-checked delivery of a byte stream. | Bi-directional connection: resends lost segments. | Endpoint also identified by (IP address, port).

→ **4.4 Protocol originators & standards**

Internet protocol (TCP/IP, etc.) originated with DARPA.

Maintained by IETF via public RFCs (Requests For Comments).

Standards can be: **De jure** – by law / of official standards bodies. | **De facto** – widely used in practice, become standards by usage.

→ **4.5 Hardware / Ethernet & MAC addresses**

+Hardware (link) layer

Implementation by **Network Interface Cards (NICs)** and drivers. | Uses **MAC (Media Access Control) addresses**, unique per interface (48-bit).

→ **Ethernet frame format (simplified)**

Preamble (8 bytes) – sync. | Destination MAC (6) | Source MAC (6) | Type (2) – indicates payload protocol (e.g. IPv4, ARP). | **Payload** – data (e.g. IP datagram, ARP message). | CRC (4) – error-detecting code.

→ **Ethernet behaviour**

Originally used CSMA/CD on shared media; nowadays mostly switched (collision-free). | Frame is best effort; errors may cause drops, not retries (reliability was sacrificed for speed).

→ **5. NETWORKING PART 2 – TCP/IP, Routing & Naming**

+Encapsulation recap

Application data → wrapped in TCP/UDP segment → wrapped in IP datagram → wrapped in Ethernet frame for transmission.

→ **5.1 IP datagram basics**

Contains an IP header + payload (usually TCP or UDP). | Header fields include: Source & destination IP addresses. | Protocol (TCP/UDP/ICMP). | Total length. | Header checksum. | TTL (Time To Live). | Flags & fragmentation info.

→ **5.2 Datagrams & fragmentation**

Each layer has a maximum payload size (MTU). | If incoming block too big, IP splits into multiple fragments (each its own datagram). | Receiver reassembles fragments before passing up.

→ **5.3 Transport layer & ports**

+Ports

16-bit number used to identify application on host. | Many applications use the network simultaneously; ports demultiplex traffic.

→ **Using ports**

Application "listens" on a port. | For TCP, app receives data from server. | Client connects to (server IP, port). | Server moves the connection to a new ephemeral port.

For UDP, app receives my data sent to my listening port.

→ **5.4 Hardware address resolution – ARP**

Software uses IP addresses, but Ethernet sends to MAC addresses. | ARP (Address Resolution Protocol) maps IP → MAC on local network. | Host A sends ARP request (broadcast at Ethernet level) asking "who has IP X?". | Host B with IP X replies with its MAC. | Mapping cached for a while, then forgotten.

→ **5.5 IP addressing & IPv6**

IPv4: 32-bit addresses, dotted decimal (e.g. 192.168.1.10). Limited space (~4.3 billion).

IPv6: 128-bit addresses, hex groups (e.g. 210E:0154:D017:9FB8). Huge space.

In address like 192.168.1.234, part identifies network (e.g. 192.168.1), rest is host interface (e.g. 234).

One computer with multiple network interfaces can have **multiple IPs**; on same network, those IPs must be different.

→ **5.6 Routing**

Routing ensures IP datagrams reach the correct network. | If destination network == current network → deliver directly to host. | Else → consult routing table.

send to gateway router / ISP. | Router has interfaces on two or more networks and forwards datagrams between them.

→ **5.7 Naming & DNS**

IP addresses are hard to remember; we use hostnames like www.singaporetech.edu.sg.

Names are hierarchical: e.g. library.singaporetech.edu.sg | sg – country (delegated to local authority). | edu – ac – sector (education). | singaportechnology – organisation. | library – specific host.

→ **DNS (Domain Name System)**

Implemented by hierarchy of DNS servers (often duplicated for reliability). | Common record types: A – name → IP address. | PTR – address → name (reverse lookup). | CNAME – alias/another name. | TXT – arbitrary text (e.g. SPF). | SOA – start of authority. | MX – mail exchanger for domain.

→ **DNS queries**

User client "talks" to transport layer (TCP/UDP). | Programmer uses sockets (same idea in C, Java, Python, etc.).

→ **Client-server model**

Server application listens on known port. | Client applications connect to server (same or different machine). | Client ↔ server communication done over TCP or UDP.

→ **Data units**

UDP = messages / datagrams (collections of bytes).

## 8. Computer Architecture: History, Structure & Organisation

→ What is inside a computer?

Blocks of electronic circuits: CPU (Central Processing Unit), | Memory (RAM, ROM, cache). | I/O controllers, buses, etc.

Connected via buses (address, data, control).

→ Computer Architecture vs Organisation

Architecture – what the computer does; functional behaviour and visible features. | Organisation – how those functions are implemented; structural relationships and hardware details. | Same architecture (e.g. ARM) can be implemented with different organisations (e.g. different pipelines, caches).

→ Layered view of computer systems

From bottom up: Electronic gates + logic circuits. | Micro-architecture & control unit. | Instruction set architecture (machine code). | Assembly language. | High-level languages (C, Java, Python). | Applications.

BIOS/firmware runs at boot to initialise hardware and load OS.

→ Von Neumann vs Harvard architectures

Von Neumann: | Single memory holds both **program instructions** and **data**. | CPU uses same bus to fetch instructions and data. | Simpler, but instruction and data fetch compete for bandwidth (bottleneck).

Harvard: Separate memories (and often buses) for **program** and **data**. | CPU can fetch instruction and data in parallel → potential speedup. | Used in many microcontrollers and DSPs.

→ CPU interaction with memory

To execute instruction: Fetch instruction from memory (often via instruction cache). | Decode instruction. | Fetch operands (from registers or memory). | Execute in ALU / other functional units. | Write results back to registers or memory.

Program Counter (PC) holds address of next instruction.

Registers store temporary data and addresses.

Instruction Cache (I-cache) stores recently used instructions near the CPU.

## 9. Information Representation & Data Formats

9.1 Number systems

→ Decimal (base 10) – digits 0–9.

→ Binary (base 2) – digits 0,1.

→ Hexadecimal (base 16) – digits 0–9, A–F.

→ Conversion notes

Binary ↔ decimal – positional weights  $2^0, 2^1, 2^2, \dots$

Binary ↔ hex – group bits in 4s.

9.2 Fixed-point integer representation

→ Unsigned n-bit integer

Range: 0...2<sup>n</sup> - 1

→ Signed two's complement

Range: -2<sup>n-1</sup> ... 2<sup>n-1</sup> - 1

To get negative: invert bits + 1.

9.3 Fixed-point fraction

Some bits represent integer, some represent fractional part.

Fraction bits correspond to 1/2, 1/4, 1/8, ...

9.4 Floating point (concept)

Number ≈ (-1)<sup>sign</sup> × 1.mantissa × 2<sup>(exponent-bias)</sup> | Allows representing a very wide range of values (very small and very large).

9.5 Binary arithmetic (ALU)

ALU performs add, subtract, multiply, divide, logic ops (AND, OR, XOR, NOT, etc.).

Addition – same rules as decimal but in base 2, with carries. | Subtraction – using borrow or by adding two's complement. | Multiplication – repeated shift and add. | Division – repeated subtract or shift and subtract. | Shifting left / right by a bits ≈  $2^k \times 2^l$ .

9.6 Memory addressing, endianness & pointers

→ Endianness

Big-endian – most significant byte at lowest address. | Little-endian – least significant byte at lowest address (e.g. x86). | Bi-endian – can be configured either way (e.g. ARM). | Programmer usually doesn't need to worry; compiler/OS handle it.

9.7 Bits, bytes, words, wordlength & effects

Bit – 0 or 1. | Nibble – 4 bits. | Byte – 8 bits. | Word – group of bits stored/addressed together; size depends on architecture (8, 16, 32, 64 bits). | Wordlength – number of bits per word. | Affects range & precision for numbers. | Larger wordlength → higher quality (e.g. audio), but more memory & bandwidth. | Trade-off: Quality vs memory size / performance.

→ Storage calculation example (audio CD)

44.1 kHz × 16 bits × 3600 s × 2 channels ×  $2.54 \times 10^8$  bits = 318 MB per channel → ~636 MB for stereo.

## 10. CPU Basics & Instruction Set Architecture (ISA)

10.1 RISC vs CISC

→ RISC (Reduced Instruction Set Computers)

Simple, fixed-format instructions (load, operate, store). | Each instruction does a small job that can usually complete in 1 clock cycle. | Pros: simple hardware, easy pipelining, fewer cycles per instruction. | Cons: more instructions per program.

→ CISC (Complex Instruction Set Computers)

Rich instructions that may perform multiple steps (load+operate+store). | Pros: fewer instructions per program; compact code. | Cons: variable length, more cycles per instruction, more complex hardware.

10.2 CPU basics

→ CPU components

Registers – small, ultra-fast storage for operands, addresses, flags. | ALU – performs arithmetic/logic. | Control unit – orchestrates fetch-decode-execute. | Internal buses – move data between units (operand buses, result bus). | Program Counter (PC) – address of next instruction. | Instruction cache / pipeline – store/fetch upcoming instructions.

→ Internal CPU bus & MAC

CPU can have multiple execution units; data moved via internal buses. | MAC (Multiply-Accumulate) unit performs multiply + add in a single step (useful in DSP).

10.3 Instruction cycle

Stages (conceptually): Fetch – get instruction from memory/I-cache using PC. | Decode – figure out what to do and which operands. | Load operands – from register or memory. | Execute – ALU / MAC / other units process. | Write back – store result, update flags, PC.

Instruction time = time to complete all stages (may be multiple clock cycles).

10.4 Instruction Set Architecture (ISA)

ISA = set of machine-level instructions & how they behave. | Includes opcodes, formats, data types, addressing modes, registers, etc. | High-level language programs are compiled to ISA instructions.

11. Memory Types, Operations & Management, Peripherals & I/O

11.1 Types of memory

→ Volatile vs non-volatile

Volatile – loses content when power is off (RAM, cache). | Non-volatile – keeps content (ROM, Flash, HDD, SSD).

→ RAM vs ROM

RAM – fast, read/write, smaller capacity, directly accessed by CPU. | ROM – slower, read-mostly, larger capacity, stores firmware/boot code.

→ DRAM vs SRAM

DRAM – stored with capacitors; needs refresh; slower but denser → main memory. | SRAM – stored with flip-flops; very fast but expensive → caches.

→ Cache & hit ratio

Registers → cache → main memory → secondary storage. | As you go down: capacity ↑, cost/bit ↓, speed ↓.

→ Cache & hit ratio

Cache stores copies of frequently used memory locations. | Cache hit – data found in cache → fast. | Cache miss – data not there → fetched from main memory, cache updated. | Performance measured by hit ratio = hits / total accesses.

11.2 Memory errors & parity

Errors can be hard (permanent) or soft (transient, e.g. noise, cosmic rays). | Parity bit: add 1 extra bit so that number of 1s is even (even parity) or odd (odd parity). | Can detect single-bit errors, but not all multi-bit errors.

11.3 Memory management unit (MMU) & virtual memory (recap)

MMU handles translation from virtual to physical addresses and implements virtual memory using paging and possibly segmentation. | Virtual memory allows program size > physical RAM by swapping pages to disk.

11.4 Peripherals & I/O

→ Peripherals

Devices that let computer interact with outside world: keyboard, mouse, monitor, printer, network, sensors/actuators, etc. | Connected via external serial/parallel ports (USB, HDMI, etc.) to I/O interfaces.

→ I/O bus and system bus

I/O bus (address, data, control) shared by multiple peripherals. | System bus = address + data + control buses connecting CPU, memory, I/O.

→ Addressing peripherals

Each device has an address (or range).

CPU sends: Address → which device. | Control signals (read/write, etc.). | Data to/from device.

11.5 System bus allocation schemes (I/O addressing)

Separate address, data, control buses for memory vs I/O. | Different instruction set for memory vs I/O. | Pros: efficient, no contention between memory and I/O. | Cons: more hardware, more complex.

Isolated I/O: Shared address & data buses; separate control lines for memory vs I/O. | Memory and I/O have different address spaces. | I/O addresses often called ports.

Memory-mapped I/O: Same address, data, control buses used for both memory and I/O. | I/O devices appear as special memory addresses. | Pros: same instructions for memory and I/O; simpler programming, simpler logic. | Cons: reduces address space available for normal memory.

12. Pipelining & Polling

12.1 Pipelining

→ Analogy (laundry)

Tasks: wash, dry, fold, put away. | Non-pipelined: finish all stages for load 1 before starting load 2. Machines idle a lot. | Pipelined: as soon as washer is free, start next load while dryer works on previous load.

→ Definition

Pipelining = overlapping execution of multiple instructions by splitting instruction execution into stages handled by different hardware units.

→ Typical CPU stages (example)

FETCH instruction. | DECODE instruction. | LOAD operands. | EXECUTE. | RESULT & SAVE.

→ Performance effect

Latency – time to complete one instruction (often slightly higher with pipeline). | Throughput – instructions completed per second (much higher once pipeline is full). | Ideal: throughput improves roughly by number of pipeline stages (if all stages same time and no hazards).

→ Hazards (only conceptually)

Structural – hardware resource conflicts. | Data – instruction depends on result of previous one. | Control – branches change instruction flow.

(You mainly need to know they reduce ideal pipeline performance.)

12.2 Polling

→ What is polling?

CPU repeatedly checks device status in a loop to see if device needs attention.

→ Pros

Very simple to program and understand. | CPU has precise control over check order.

→ Cons / trade-offs

Wastes CPU time checking when nothing has changed. | Latency between checks can be large (especially with multiple devices).

→ Multiple-device polling

Main loop checks each device in turn (GPIOs, timers, etc.). | Larger number of devices → bigger latency before some device is serviced.

## 13. Interrupts & Direct Memory Access (DMA)

13.1 Interrupts

→ Idea

Interrupt = hardware or software signal that causes CPU to temporarily stop current code and run an **Interrupt Service Routine (ISR)**.

Device can only CPU only when it needs attention. | CPU doesn't waste time polling. | Allows CPU to continue other work until event occurs.

→ Types

Hardware interrupt – from external devices (keyboard, mouse, timers, network, etc.). | Software interrupts / exceptions – triggered by CPU (e.g. divide-by-zero, system call instruction).

→ Interrupt vector table

Fixed memory table mapping **interrupt numbers** → **ISR addresses**. | Each interrupt has an entry.

→ What happens on an interrupt?

CPU detects interrupt, suspends current execution. | CPU saves context (PC, registers, flags). | CPU looks up ISR address from **vector table**. | CPU jumps to ISR and executes it (service device, clear flags). | ISR returns; CPU restores saved context. | Execution resumes where it left off.

→ Pros / cons

Efficient CPU usage; good responsiveness.

→ Code more complex; must handle concurrency and re-entrancy.

13.2 DMA – Direct Memory Access

→ Motivation

Many I/O operations involve moving large blocks of data between device and memory without involving CPU for each byte/word. | CPU sets up DMA: source address, destination address, length, direction. | DMA then performs transfer over system bus. | When done, DMA triggers an interrupt to notify CPU.

→ Benefits

Allows CPU to do other work while data moves. | Higher throughput for bulk transfers (disk, network, etc.).

## LAB 1 — Linux + Shell + C (VERY EXAM-POPULAR)

### Essential Commands

Directory ops: mkdir, cd, pwd, ls, ls -la, file \* | Viewing files: cat, head, tail | Counting: wc – lines, words, bytes | Redirection: > overwrite, >> append | Pipes: cmd1 | cmd2 | Search/transform: grep, grep -E, sed 's/a/b/g' | History: history, man cmd, cmd -help

### Regex Quick Patterns

Contains "zoo": grep -E 'zoo' | "zoo+" = "zoo", "zooo..." | "zo+" = "zo", "zoo..." | Starts with: ^zoo | Word boundary: \zoo | Seven-letter words c....n: grep -E '^c....n\$'

### Loops, If, Variables

For loop: for ((i=1;i<=10;i++)); do echo \$i; done

Arithmetic: z=\$((x+y))

If (strings): [ \$S1 == \$S2 ]

If (numbers): [ \$S1 -eq -ne -lt -le -ge ]

### Wildcards

\* = match all files Example: /etc/\*.conf

### Shell script basics

#!/bin/bash echo "one"

Make executable: chmod u+x rxfile.sh

Run: ./file.sh

### C compile

Compile: gcc -o myprog myprog.c | Redirect output: ./myprog > dump

### Find user / passwd line

Locate user in passwd: grep '^www:www:' /etc/passwd

### Web log analysis (very examinable)

Unique hosts: awk '{print \$1}' access.log | sort -u | wc -l

Count IPs starting with 148: awk '{print \$1}' access.log | grep '^148.' | wc -l

Most accessed file: awk '{print \$1}' access.log | sort | uniq -c | sort -h

## LAB 2 – Networking + Sockets

### Network Info Commands

Interfaces:

ifconfig or nmcli d show | grep interface

MAC address: ip link | IP (v4/v6): ip addr | ARP table: cat /proc/net/arp | Device stats: cat /proc/net/dev

### /proc system info

Memory: cat /proc/meminfo | CPU: cat /proc/cpuinfo | Disk: cat /proc/diskstats | Per-process info: /proc/<PID>/status (VmSize, VmPeak, etc.)

### Connectivity

Check interface up: ifconfig | nmcli g | Ping IP: ping 8.8.8.8 | Ping domain: ping bcc.co.uk | DNS queries: nslookup, dig | Trace route: tracepath host or traceroute host

### Client / Server architecture (C TCP sockets)

Client uses:

gethostbyname() | socket(AF\_INET, SOCK\_STREAM, 0) | connect() | write()

Server uses:

socket() | bind() | listen() | accept() | read() | write()

Compile:

gcc client.c -o client | gcc server.c -o server

Run:

Server: ./server <port> | Test via: telnet localhost <port>

### Manual HTTP (EXAM TRICK)

telnet <host> 80 | GET /HTTP/1.1 HOST: <host>

## LAB 3 – Raspberry Pi + GPIO + Ultrasonic Sensor + Python

### Playing Sound (pygame)

import pygame

pygame.init()

snd.play()

### gpiozero Buttons + LEDs

from gpiozero import Button, LED

bttn = Button(17)

led = LED(4)

bttn.when\_pressed = led.on

bttn.when\_released = led.off

Blink:

while True:

led.on

time.sleep(1)

led.off()
 time.sleep(1)

### Button + Sound:

bttn.when\_pressed = snd.play

### Ultrasonic Sensor (HC-SR04)

(/MOST EXAMINABLE SECTION)

### Distance formula

Where:

speed of sound = 343 m/s = 34300 cm/s

t = pulse duration (seconds)

### Key Measurements

Trigger pin HIGH for 10 us

Echo pin HIGH for travel-time

Min range 2 cm, max – 4 m

### Essential Code Pattern

GPIO.setmode(GPIO.BCM)

TRIG=23; ECHO=24

GPIO.setup(TRIG, GPIO.OUT)

GPIO.setup(ECHO, GPIO.IN)

GPIO.output(TRIG, False)

time.sleep(2)

# Trigger

GPIO.output(TRIG, True)

time.sleep(0.00001)

GPIO.output(TRIG, False)

# Measure pulse

while GPIO.input(ECHO)==0:

start = time.time()

while GPIO.input(ECHO)==1:

end = time.time()

distance = (end-start)\*17150

distance = round(distance,1)

### Beep Frequency Logic

(Assignment)

def beep(freq):

d = distance

if d > 50: return -1 # no beep

elif d >= 30: return 1 # 1 beep/sec

elif d >= 20: return 0.5 # 2 beeps/sec

elif d >= 10: return 0.25 # 4 beeps/sec

else: return 0 # continuous

Main loop:

freq = beep( freq0 )

if freq == -1: off